

Comments

```
// line: from double-slash until the end of the line
/* block: from open to close; multi-line & nested. */
/// line documentation for functions, structs, etc.
//! line documentation for crates and modules.
```

Variables

Binding names to values

- The "let" keyword binds a variable name to a value. (The variable becomes the owner of the value). The drop function destroys a value.
- A variable name can be bound to a different value – shadows the old value for the rest of the code-block.

```
let x = "outer"; // BINDING
{
    // new code-block
    let x = "inner"; // NEW BINDING, outer x shadowed
    let x = "inner2"; // NEW BINDING, inner x shadowed
} // end block, both inner x dropped
println!("{}", x); // prints "outer"
```

- Variables are immutable by default. The "mut" keyword is required for mutable variables.

```
let mut y = 5; // Mutable binding
y = 10; // Assignment: Same binding, new value
```

- The variable type must be known at compile time, but this can often be inferred by the compiler (as above).

Static

- static defines a global variable with a fixed memory location. It has a 'static lifetime, which means the data lives for the entire duration of the program.

```
static HELLO: &str = "Hello, world!"; // &'static str
```

- Static variables must be declared at the global (module) level, not inside functions or blocks.

Const

- const in Rust creates compile-time constants (which are not a variable) – the const value is computed during compilation and embedded directly into your code. A const does not reside as a variable in memory. Constants must be typed.

```
const PI_OVER_2: f32 = 3.14159265 / 2.0f32;
```

Types

Scalar types (stored on the stack)

- Scalar types are integers (either signed or unsigned), floating point numbers, characters and Booleans.
- The number after the numeric types indicates how many bits are used to represent the number. Note: isize and usize are pointer-sized integer types whose size depends on the target machine architecture.
- If not explicitly stated, the default integer type is i32. The default floating type is f64.

```
let a: i32 = -42; // i8, i16, i32, i64, i128, isize
let b: u32 = 42; // u8, u16, u32, u64, u128, usize
let c: f64 = 3.14; // f32, f64
let d: bool = true; // Boolean
let e: char = '★'; // Unicode character, single quotes
```

Built-in (AKA primitive) compound types

- Tuples:** fixed length collection of values (can be of different types) – typically stored on the stack.

```
let tup: (i32, f64, u8) = (500, 6.4, 1);
let (x, y, z) = tup; // destructuring
let one = tup.2; // index from zero accessible
let empty = (); // unit or empty tuple
```

- Arrays:** fixed length collection of values of the same type – typically stored on the stack.

```
let a = [1, 2, 3, 4, 5];
let months = ["January", "February", "March"];
let b: [i32; 5] = [1, 2, 3, 4, 5];
let repeated = [3; 5]; // [3, 3, 3, 3, 3]
let second_value = b[1]; // indexed from 0
```

- A **tuple struct** is a struct with fields that are accessed by position (like a tuple) rather than name. Values can be of a different type.

```
struct Empty()
struct UserID(i32)

struct Point(i32, i32); // 2D point
let point = Point(10, 20);
let x = point.0; // Access first field
let y = point.1; // Access second field
```

Nominal (named) types; AKA abstract data types

- Struct:** Like tuples, but with named fields – can encapsulate code as well as data using the impl keyword (discussed later). With encapsulated code, somewhat like classes in C++ (without inheritance).

```
struct Member {
    name: String,
    age: u32,
    is_active: bool,
}

let member_1 = Member {
    name: String::from("Alice Smith"),
    age: 31,
    is_active: true,
};

let member_2 = Member {
    name: String::from("Bob Johnson"),
    is_active: false,
    ..member_1 // copy member_1 for remaining fields
};

let name = member_1.name;
```

- Enum:** one of multiple alternative data items/values. Unlike enums in other languages, Rust enums can hold data values. Enums can have code impl blocks just like structs (discussed later).

```
// enums without data values
enum Colour {
    Red,
    Green,
    Blue,
}
let favourite_colour = Colour::Red;

// enums with data values
enum Shape {
    Circle(f64), // radius
    Rectangle(f64, f64), // width, height
    Triangle(f64, f64), // base, height
}
let my_shape = Shape::Circle(2.5);
```

Strings (introduction)

- String slices (type: `&str` – a borrowed string slice) – immutable by default – double quotes – fixed size. Like a pointer to a slice of Unicode characters from either a "literal string" known at compile time, or the contents of a String struct.

```
let greeting = "Hello, world!"; // Type: &str
let multiline = "This is a
multiline string";
```

- String – a built-in struct that manages mutable, ownable, variable length strings (stored on the heap).

```
// Creation
let empty = String::new();
let from_string_literal = String::from("hello");
let pre_alloc_capacity = String::with_capacity(100);
let conversion = "hello".to_string();

// Building
let mut s = String::from("Hello");
s.push('!'); // Add single character
s.push_str(" World!"); // Add string slice
s.insert(5, ','); // Insert char at index
s.insert_str(6, " beautiful"); // Insert string at index
s.remove(0); // Remove char at index

// Inspection
let s = String::from("Hello, 世界!"); // Hello world
let length_in_bytes = s.len();
let character_count = s.chars().count();
let is_empty = s.is_empty();
let current_capacity = s.capacity();

// Conversion
let as_string_slice = s.as_str();
let as_byte_slice = s.as_bytes();

// Case
let a = s.to_lowercase(); // returns new String
let a = s.to_uppercase(); // returns new String

// Trimming white space
let a = s.trim(); // both ends
let a = s.trim_start(); // beginning only
let a = s.trim_end(); // end only
let a = s.trim_matches(|c| c == ' '); // custom
let a = s.trim_start_matches("pre"); // remove prefix
let a = s.trim_end_matches("suf"); // remove suffix

// Searching and replacing
let b = s.contains("sub"); // substring exists
let b = s.starts_with("pre"); // prefix check
let b = s.ends_with("suf"); // suffix check
let position = s.find("pattern"); // first occurrence
let position = s.rfind("pattern"); // last occurrence
let a = s.replace("old", "new"); // replace all
let a = s.replacen("old", "new", 2); // first n only

// split - creates an iterator – by char ('c') or "str"
// Typically needs to be collected into a Vector
let text = "apple,banana,cherry";
let fruits: Vec<&str> = text.split(',').collect();
println!("Fruits: {}", fruits.join(", "));

let text = " Hello world! ";
let w: Vec<&str> = text.split_whitespace().collect();
println!("{}", w); // Output: ["Hello", "world!"]

// Slicing strings
let (s, start, end) = ("01234567890", 3, 7);
let o = s.get(start..end); // safe slice, returns Option

// Iteration over characters
let text = "Hello 世界";
for ch in text.chars() {
    println!("{}", ch);
}
```

Data containers

- Rust has many useful data containers (that are built-in structs with methods), including dynamic arrays, hash-maps and hash-sets (similar to lists, dictionaries and sets in Python).

- Vectors – dynamic arrays

```
let mut vec: Vec<i32> = Vec::new();
let mut vec: Vec<String> = Vec::with_capacity(10);
let mut vec = vec![1, 2, 3]; // a useful macro

vec.push(4); // add to end
let item = vec.pop(); // take from end, returns Option
vec.insert(0, 0); // insert - (index, item)
vec.remove(0); // remove at index
let len = vec.len();
let item = vec[0]; // direct access (can panic)
let item = vec.get(0); // safe access, returns Option
vec.clear(); // remove all items

for item in &vec { println!("{}", item); } // borrow
for item in vec.iter_mut() { *item += 1; } // m borrow
for (i, e) in vec.iter().enumerate() {
    println!("{}", i, e);
}
for item in vec { println!("{}", item); } // consume
```

- Hash Maps

```
use std::collections::HashMap; // import

// Think about borrowed vs owned keys and values
let mut map: HashMap<&str, &i32> = HashMap::new();
let mut map: HashMap<String, i32> = HashMap::new();
let mut map = HashMap::from([("key", "value")]);

map.insert("key", "value");
let value = map.get("key"); // returns Option<&V>
let value = map["key"]; // direct access (can panic)
map.remove("key");
let contains = map.contains_key("key");

map.entry("key").or_insert("val"); // insert if missing

// a counter
let mut cnt: HashMap<&str, i32> = HashMap::new();
cnt.entry("key").and_modify(|v| *v += 1).or_insert(1);

for (key, value) in &map {
    println!("{}", key, value);
}
for key in map.keys() { println!("{}", key); }
for value in map.values() { println!("{}", value); }
```

- Hash Sets

```
use std::collections::HashSet;

let mut set: HashSet<i32> = HashSet::new();
let mut set = HashSet::from([1, 2, 3]);

set.insert(4);
set.remove(&2);
let contains = set.contains(&3);
let len = set.len();

let set1 = HashSet::from([1, 2, 3]);
let set2 = HashSet::from([3, 4, 5]);

let union: HashSet<_> =
    set1.union(&set2).collect();
let intersection: HashSet<_> =
    set1.intersection(&set2).collect();
let difference: HashSet<_> =
    set1.difference(&set2).collect();

for item in &set { println!("{}", item); }
for (i, e) in set.iter().enumerate() {
    println!("{}", i, e);
}
for item in set { println!("{}", item); }
```

Ownership and Borrowing (Pointers in Rust)

Core Ownership Rules

- Values only have one variable owner at any given time
- If the owner goes out of scope, the value is dropped. Values can be dropped at any time manually.

```
let x = String::new(); // x owns the empty String value
drop(x);               // x is inaccessible after this line
```

- Values also go out of scope when the code execution leaves the code block {} in which they were declared.

Ownership Transfer (Moving)

- With non-Copy types, ownership moves when you assign it or pass it to a function by value.
- Non-Copy types (eg. String, Vec, HashMap, custom structs as well as mutable references) are always moved, invalidating the original variable.

```
let v1 = vec![1, 2, 3];
let v2 = v1; // v1 is moved, only v2 is valid

let mut s = String::from("Hello sweetie");
let r1 = &mut s; // take a mutable reference to s
let r2 = r1;     // r1 is moved to r2, only r2 is valid

fn takes_ownership(s: String) {} //s dropped at close
fn main() {
    let mine = String::new();
    takes_ownership(mine); // ownership of mine moved
    // "mine" is no longer valid
}
```

- Note: rust has a "deep copy" .clone() method that many structs implement. It produces a new independent value, but it must be called explicitly.

```
let v = vec![1, 2, 3];
let w = v.clone(); // w is copied, both v and w valid
```

- Caution: cloning large data structures or repeated cloning in a loop can be time and space inefficient.
- Mutable references cannot be copied or cloned.

- Copy types: Values that implement the .copy() trait are copied and not moved. Note: primitive scalar types and immutable references are always copied.

```
let x = 5; // int, float, bool and char are primitive
let y = x; // x is copied, both x and y are valid

let s = String::from("Hello sweetie");
let r1 = &s; // take an immutable reference to s
let r2 = r1; // r1 is copied, now two references to s
```

Borrowing Rules (apply to reference types &T, &mut T)

- Borrowing in Rust is like taking a reference in C/C++. A reference is a smart pointer to a value.
- **Immutable Borrowing (&T)**
 - Multiple immutable references simultaneously
 - Can read but cannot modify
 - Original owner retains ownership

```
let s = String::from("hello");
let r1 = &s; // immutable borrow
let r2 = &s; // multiple immutable borrows OK
println!("{}", r1, r2); // both valid
```

- While an immutable borrow is active, the owner cannot: mutate the value, move the value, nor create mutable borrows. The owner can read the value.

```
let mut x = vec![1, 2, 3];
let last = x.last().unwrap(); // immutable borrow &x[2]
x.push(4); // FORBIDDEN - owner cannot mutate x
println!("{}", last); // Rust cannot certify last ref
```

Mutable Borrowing (&mut T)

- Only ONE mutable reference allowed at a time
- Cannot have immutable references while a mutable reference exists
- Can read and modify

```
let mut s = String::from("hello");
let r1 = &mut s; // mutable borrow
r1.push_str(" world");
```

- While a mutable borrow is active, the owner cannot: read the value, mutate the value, move the value, nor create any other borrows (mutable or not).

Automatic dereferencing of borrowed items

- Like C/C++, Rust uses "*" as the dereference operator.
- Unlike C/C++ (which uses "->"), Rust uses "." as the dereference for borrowed struct member values.
- Rust doesn't need the explicit "*" dereference in many situations.

```
let owned = String::from("Hello friend");
let b = &owned; // immutable borrow

// These all work without explicit dereferencing:
println!("{}", b.len()); // Method calls auto-derefs
println!("{}", b);       // Display auto-derefs
let first_char = b[0..1]; // Indexing auto-dereferences
```

- But there are some situations where the explicit dereference is required:
 - Assignment through a reference

```
let mut owned = 5;
let mut borrowed = &mut owned;
*mut_borrowed = 10; // Explicit dereference to assign
```

- Comparison of referenced values

```
let owned = 5;
let borrow = &owned;

// Compare the value, not the reference
if *borrow == 5 { println!("Equal!"); }
```

Common Patterns with ownership/borrowing

- Function parameters – prefer borrow to move:

```
fn process(s: &String) { } // borrowing (preferred)
fn take_ownership(s: String) { } // moving
```

- Returning references:

```
fn get_first(v: &Vec<i32>) -> &i32 {
    &v[0] // lifetime tied to input parameter
}
```

- Split borrows:

```
let mut v = vec![1, 2, 3];
let (first, rest) = v.split_at_mut(1);
// first references &mut [1], rest refs &mut [2, 3]
```

(Reference) Lifetimes

What are lifetimes?

Lifetimes are Rust's way of tracking how long references are valid. Every reference has a lifetime, even if not explicitly written. The compiler ensures references don't outlive the data they point to.

```
&'a T // Reference with lifetime 'a
&'a mut T // Mutable reference with lifetime 'a
```

Function Lifetimes

When a function returns a reference, you must specify how input and output lifetimes relate:

```
// The returned reference lives as long as both inputs
fn longest<'a>(x: &'a str, y: &'a str) -> &'a str {
    if x.len() > y.len() { x } else { y }
}

// Different lifetimes - tied only to first param
fn first_word<'a, 'b>(s: &'a str, _other: &'b str)
-> &'a str {
    s.split_whitespace().next().unwrap_or("")
}
```

Lifetime Elision Rules (when you can omit them)

The compiler automatically infers lifetimes that are obvious from the context, where:

1. Each input reference gets its own lifetime
2. If there's one input lifetime, output gets the same lifetime
3. If there's &self or &mut self, output gets self's lifetime

```
fn print(s: &str) // Elided
fn print<'a>(s: &'a str) // Explicit equivalent

fn get_first(s: &str) -> &str // Elided
fn get_first<'a>(s: &'a str) -> &'a str // Explicit

impl Thing {
    fn get_name(&self) -> &str // Elided (rule 3)
    fn get_name<'a>(&'a self) -> &'a str // Explicit
}
```

Struct Lifetimes

Structs storing references need lifetime parameters:

```
struct Book<'a> {
    // These references must live as long as Book
    title: &'a str,
    auth: &'a str,
}

impl<'a> Book<'a> {
    fn new(title: &'a str, auth: &'a str) -> Book<'a> {
        Book { title, auth }
    }
}

let title = String::from("Rust");
let auth = String::from("Klabnik");
let book = Book::new(&title, &auth);
// book cannot outlive title or author
```

Static Lifetime

'static means the reference lives for the entire program:

```
let s: &'static str = "literals have static lifetime";

// T: 'static - T contains no non-static references
fn needs_static<T: 'static>(t: T) {}
```

Common Patterns

```
// Multiple lifetimes with constraints
fn h<'a, 'b>(x: &'a str, y: &'b str) -> &'a str {
    // 'b: 'a means 'b lives at least as long as 'a
    if x.len() > 5 { x } else { y }
}

// Lifetime in where clause
fn where_example<'a, T>(item: &'a T) -> &'a T
where
    T: std::fmt::Display + 'a
{
    println!("{}", item);
    item
}
```

Operators

Arithmetic operators

```
let a = 10;
let b = 3;

// Basic arithmetic
let sum = a + b; // 13
let difference = a - b; // 7
let product = a * b; // 30
let quotient = a / b; // 3 (integer division)
let remainder = a % b; // 1

// Floating point division
let (x, y) = (10.0, 3.0);
let float_quotient = x / y; // 3.333...

// Unary operators
let negative = -5;
```

- Note: there are no implicit type conversions for numeric values – you must explicitly cast types to the same type for arithmetic operations.

```
// Numeric conversions
let x = 42i32; // i32
let y = x as f64; // Cast to f64
let z = y as u32; // Cast to u32
let byte = z as u8; // Cast to u8
```

- But Rust can use context to infer the types of literals.

```
let x: f32 = 3.14; // 3.14 is inferred to be f32
let y = x + 2.0; // 2.0 is inferred as f32 to match x
```

- Note: the std crate includes many maths functions:
 - Integer methods – abs(), pow(), etc.
 - Floating methods – sqrt(), sin(), exp(), ln(), etc.
 - Min/Max/Rounding – min(), max(), round(), etc.

Assignment operators

```
let mut x = 5;

// Basic assignment
x = 10;

// Compound assignment operators
x += 3; // x = x + 3
x -= 2; // x = x - 2
x *= 4; // x = x * 4
x /= 2; // x = x / 2
x %= 3; // x = x % 3

// Bitwise compound assignment
x &= 0b1010; // x = x & 0b1010
x |= 0b0101; // x = x | 0b0101
x ^= 0b1111; // x = x ^ 0b1111
x <<= 2; // x = x << 2
x >>= 1; // x = x >> 1
```

Comparison operators

```
let (a, b) = (5, 10);

// Equality and inequality
let equal = a == b; // false
let not_equal = a != b; // true

// Ordering
let less = a < b; // true
let greater = a > b; // false
let less_equal = a <= b; // true
let greater_equal = a >= b; // false

// String comparison
let s1 = "apple";
let s2 = "banana";
let str_less = s1 < s2; // true (lexicographical)
```

Logical Operators

```
let a = true;
let b = false;

// Logical AND (short-circuiting)
let and_result = a && b; // false

// Logical OR (short-circuiting)
let or_result = a || b; // true

// Logical NOT
let not_a = !a; // false
let not_b = !b; // true
```

- In Rust, the logical operators && (and) and || (or) are lazy (short-circuited). That means they only evaluate the right-hand side if they really need to (same as C/C++).

Bitwise operators

```
let a = 0b1010; // 10 in binary
let b = 0b1100; // 12 in binary

// Bitwise AND
let and = a & b; // 0b1000 (8)

// Bitwise OR
let or = a | b; // 0b1110 (14)

// Bitwise XOR
let xor = a ^ b; // 0b0110 (6)

// Bitwise NOT
let not_a = !a; // Flips all bits

// Bit shifting
let left_shift = a << 2; // 0b101000 (40)
let right_shift = a >> 1; // 0b0101 (5)
```

Range operators

```
// Inclusive range (includes end)
let inclusive = 1..=5; // 1, 2, 3, 4, 5

// Exclusive range (excludes end)
let exclusive = 1..5; // 1, 2, 3, 4

// Usage in loops
for i in 1..=3 {
    println!("{}", i); // Prints 1, 2, 3
}

// Slicing
let arr = [1, 2, 3, 4, 5];
let slice = &arr[1..4]; // [2, 3, 4]
let slice2 = &arr[..3]; // [1, 2, 3] (from start)
let slice3 = &arr[2..]; // [3, 4, 5] (to end)
let slice4 = &arr[..]; // [1, 2, 3, 4, 5] (all)
```

Reference and dereference operators (&, *)

```
let x = 5;
let r = &x; // Create a reference to x
let val = *r; // Dereference: get the value r points to

let mut y = 10;
let mr = &mut y; // Mutable reference
*mr = 20; // Dereference and assign
```

Index operator – []

```
let mut arr = [1, 2, 3, 4, 5];
let first = arr[0]; // 1 - copied assignment
arr[0] = 100; // Replace first value
```

Field access operator - .

```
struct Point { x: f64, y: f64 }
let point = Point { x: 5, y: 10 };
let x_coord = point.x; // Field access - dot operator
```

Flow Control

Statements, expressions and code blocks

- Statements** in Rust are semi-colon terminated.
- Expressions**, evaluate to a value, and do not have semi-colon terminators.
- A **code block** is just a chunk of code wrapped in curly braces {}. Code blocks can return a value, if the last element in the code block is an expression. If the last element is a statement, the unit tuple is returned ().

```
let y = {
    let a = 2;
    let b = 3;
    a + b // no semicolon, value of block is returned
};
```

- Code blocks are important in Rust because:
 - They control scope – contain temporary variables.
 - They return values – like inline mini-functions.
 - Resource management – when a block ends, local variables are dropped automatically.

Functions

- Functions are a key unit of flow control. Every program starts with a "main" function.

```
fn greet(name: &str) {
    println!("Hello, {}!", name);
}

fn main() {
    greet("Alice"); // calling the function
}
```

- In idiomatic Rust, functions that return a value, are written as a code block that ends with an expression.
- A return statement exists and is used for the early exit from a function.

```
fn check(n: i32) -> &'static str {
    if n < 0 {
        return "negative";
    }
    "non-negative" // No semicolon! Returned expression
}
```

- Arguments are passed by value or by reference.

```
fn pass_by_value(n: i32) {} // takes a value
fn pass_by_reference(&n: i32) {} // accepts a reference
fn pass_mut_ref(&mut n: i32) {} // accepts a mut ref
```

Macros

- Rust has built-in macros (that look like function calls, but with a name ending with a bang! – see println! above). Macros are expanded into Rust code at compile time. Note: you can write your own macros.
- Commonly used built-in macros include:

```
// String and Vector creation
let numbers = vec![1, 2, 3]; // create a Vector
let s = format!("{}", 2, 3, 5); // Strings

// Output
println!("Hello, {}!", "world");
print!("Hello, {}!", "world"); // without a newline
dbg!(2 + 2); // provide a debug print with line number

// Stop running if something is unexpected
assert_eq!(2 + 2, 4);
assert!(true);
panic!("Something went wrong!");
```

- Writing macros is not covered by this cheat sheet.

Raising an error

- There are no exceptions in Rust. If a function can fail, typically it returns an `Option` or `Result` enum.

```
fn divide(x: i32, y: i32) -> Result<f64, String> {
    if y == 0 {
        Err("Cannot divide by zero".to_string())
    } else {
        Ok(x as f64 / y as f64)
    }
}

fn divide2(x: i32, y: i32) -> Option<f64> {
    if y == 0 {
        None
    } else {
        Some(x as f64 / y as f64)
    }
}
```

Handling errors

- .unwrap() or .expect()** methods: These cause your program to panic if an error is detected. Note: Useful tools when prototyping, best not used in production.

```
let number: i32 = "42".parse()
    .expect("Failed to parse the number");
println!("The number is {}", number); // will be 42
```

- The `unwrap` method does not provide a message.

```
let number: i32 = "42".parse().unwrap();
println!("The number is {}", number); // will be 42
```

- .unwrap_or()** can be used to provide a default.

```
let number: i32 = "abc".parse().unwrap_or(0);
println!("Parsed number = {}", number); // will be 0
```

- .unwrap_or_else()** allows you to compute the default dynamically. This often done in a closure (see below).

```
let number: i32 = "abc".parse().unwrap_or_else(|| {
    println!("Value was missing, using default");
    42
});
println!("Parsed number = {}", number); // will be 42
```

- ? (error propagation) operator:** the `?` operator unwraps the value. If `None/Err` it returns early.

```
use std::fs;
fn main() -> std::io::Result<()> {
    let contents = fs::read_to_string("hello.txt")?;
    println!("{}", contents);
    Ok(())
}
```

- ? (error propagation) operator with From:** The `From` implementations tell Rust how to convert other error types into your error type, making `?` work seamlessly.

```
#[derive(Debug)]
enum MyError {
    Io(std::io::Error),
    Parse(std::num::ParseIntError),
}

impl From<std::io::Error> for MyError {
    fn from(err: std::io::Error) -> Self {
        MyError::Io(err)
    }
}

impl From<std::num::ParseIntError> for MyError {
    fn from(err: std::num::ParseIntError) -> Self {
        MyError::Parse(err)
    }
}

// Now ? automatically converts errors to MyError
fn read_number() -> Result<i32, MyError> {
    let text = std::fs::read_to_string("file.txt")?;
    let num = text.trim().parse()?;
    Ok(num)
}
```

- Repackaging:** Sometimes you will need to repackage an error to the correct type (with the `.map_err()` method and an appropriate closure) for the current function before propagation.

```
use std::fs;
use std::io;

fn read_number_from_file(path: &str) -> io::Result<i32> {
    let number = fs::read_to_string(path)?
        .trim()
        .parse()
        .map_err(|e|
            io::Error::new(io::ErrorKind::InvalidData, e)
        );
    Ok(number)
}

fn main() -> io::Result<()> {
    let n = read_number_from_file("number.txt");
    println!("Number: {}", n);
    Ok(())
}
```

- Match (see below)** can be used to handle errors:

```
use std::fs::File;
use std::io::{self, Read};

fn main() {
    let result = File::open("hello.txt");

    match result {
        Ok(mut file) => {
            let mut contents = String::new();
            file.read_to_string(&mut contents)
                .expect("Failed to read file.");
            println!("File contents:\n{}", contents);
        }
        Err(e) => {
            println!("Failed to open file: {}", e);
        }
    }
}
```

If expressions (note: no semicolons inside the curly {})

```
let num = 5;
let r = if num > 0 {"positive"} else {"not positive"};
// Note all expressions must return the same type
```

If statements

```
if number < 0 {
    println!("Negative");
} else if number == 0 {
    println!("Zero");
} else if number < 10 {
    println!("Small positive");
} else {
    println!("Large positive");
}
```

While loops

```
let mut n = 0;
while n < 10 {
    if n == 5 {
        break; // exit the loop completely
    }
    if n % 2 == 0 {
        n += 1;
        continue; // skip printing even numbers
    }
    println!("{}", n);
    n += 1;
}
println!("Loop stopped at n = {}", n);
```

Loop until you break

```
let mut count = 0;
loop {
  println!("count = {}", count);
  count += 1;
  if count == 3 {
    break; // exit loop
  }
}
```

- You can break with a value

```
let result = loop {
  let x = 2 + 2;
  break x; // returns 4
};
```

For loops over things that are iterable

```
let numbers = [10, 20, 30];
for n in numbers {
  println!("n = {}", n);
}
```

- Also, you can loop over ranges

```
for i in 0..5 { // 0,1,2,3,4
  println!("{}", i);
}
```

Match – must cover all possible cases

```
let n = 2;
match n {
  1 => println!("one"),
  2 | 3 => println!("two or three"), // or
  4..6 => println!("between 4 and 6"), // ranges
  n @ 7..9 => println!("{}", n), // @ binding
  _ => println!("something else"), // default
}
```

- Match guards and conditional matching

```
// match
let result = match n {
  x if x % 2 == 0 => "even", // conditional match
  _ => "odd",
};
println!("{}", n, result);
```

- Tuple destructuring in a match

```
let pair = (0, -2);
match pair {
  (0, y) => println!("First is 0, second = {}", y),
  (x, 0) => println!("Second is 0, first = {}", x),
  _ => println!("No zeros"),
}
```

- Matching enums

```
enum Shape {
  Circle(f64),
  Rectangle(f64, f64),
}
let shape = Shape::Circle(2.0);

let (name, area) = match shape {
  Shape::Circle(r) => ("Circle", 3.14159265 * r * r),
  Shape::Rectangle(w, h) => ("Rectangle", w * h),
};
println!("{}", shape, "has area of {} sq units.", name, area);
```

if let – match on a single value

```
let some_value = Some(42);
if let Some(x) = some_value { // Will not match None
  println!("Got value: {}", x);
}
```

While let – loop while matching

```
let mut stack = vec![1, 2, 3, 4, 5];

// Keep popping until the vector is empty
while let Some(value) = stack.pop() {
  println!("Popped: {}", value);
}
```

Closures – anonymous functions

- Closures are anonymous functions that can capture variables from their surrounding environment.

```
let add = |x, y| x + y;
println!("{}", add(5, 3)); // 8
```

- Or...

```
let (x, y) = (5, 8);
println!("{}", (|a, b| a + b)(x, y)); // borrows
println!("{}", (|a, b| a + b)(x, y)); // consumes
```

- Many methods take closures as an argument

```
let numbers = vec![1, 2, 3, 4, 5, 6];
let even: Vec<i32> = numbers
  .into_iter() // consuming iterator
  .filter(|x| x % 2 == 0) // keep evens
  .collect(); // collect in a vector
```

Iterators (often better than for loops)

- Creating iterators

```
let v = vec![1, 2, 3];
v.iter() // Iterator over &T (borrows)
v.iter_mut() // Iterator over &mut T (mut borrows)
v.into_iter() // Iterator over T (takes ownership)

// From ranges
(0..10) // 0 to 9
(0..=10) // 0 to 10 inclusive

// From arrays/slices
[1, 2, 3].iter()
"hello".chars() // Iterator over characters
"hello".bytes() // Iterator over bytes
```

- Consuming iterators

```
// Collectors
iter.collect::<Vec<_>>() // Collect into Vec
iter.collect::<HashSet<_>>() // Collect into HashSet

// Single value consumers
iter.sum::<i32>() // Sum all elements
iter.count() // Count elements
iter.last() // Get last element
iter.nth(5) // Get element at index 5
iter.find(|x| x > &5) // Find first matching
iter.max() / iter.min() // Maximum/minimum

// Boolean consumers
iter.all(|x| x > 0) // True if all match
iter.any(|x| x > 0) // True if any match

Iterator adapters (chainable)
// Transforming
iter.map(|x| x * 2) // Transform each element
iter.filter(|x| x % 2 == 0) // Keep matching elements

// Taking/skipping
iter.take(5) // Take first 5 elements
iter.skip(5) // Skip first 5 elements
iter.take_while(|x| x < &10) // Take while condition
iter.skip_while(|x| x < &10) // Skip while condition

// Combining
iter.chain(other_iter) // Concatenate iterators
iter.zip(other_iter) // Pair up elements: (a, b)
iter.enumerate() // Add index: (index, value)

// Folding – reduce to a single value
iter.fold(0, |acc, x| acc + x) // An accumulator
```

```
// Scanning - result for each step in the process
iter.scan(0, |state, x| {           // Cumulative sum
    *state += x;
    Some(*state)
})

// Reversing (if double-ended)
iter.rev()
```

Common patterns

```
// Process and collect
let results: Result<Vec<_>, _> =
    items.iter()
        .map(|x| process(x)) // .map(process)
        .collect();          // Stops on first Err

// Window/chunk operations
v.windows(2) // Sliding window: [[1,2], [2,3], ...]
v.chunks(2)  // Non-overlapping: [[1,2], [3,4], ...]

// Partition into two collections
let (evens, odds): (Vec<_>, Vec<_>) =
    nums.iter().partition(|x| x % 2 == 0);

// Custom iterator loop with mutable state
let mut iter = vec![1, 2, 3].into_iter();
while let Some(x) = iter.next() {
    println!("{}", x);
}
```

Implementation Blocks (impl)

- impl blocks let you define methods and associated functions for types (often structs and enums, but they are not limited to structs and enums).

```
struct Rectangle {
    width: f64,
    height: f64,
}

impl Rectangle {
    // Associated function (called with ::)
    fn new(width: f64, height: f64) -> Rectangle {
        Rectangle { width, height }
    }

    fn area(&self) -> f64 { // Method (called with .)
        self.width * self.height
    }

    fn destroy(self) { // Method that consumes self
        println!("Rectangle destroyed!");
        // self is moved and dropped
    }
}
```

- You can have multiple implementation blocks for the same struct or enum.

```
impl Rectangle {
    fn perimeter(&self) -> f64 {
        2.0 * (self.width + self.height)
    }
}
```

Traits

- Traits are guarantees for shared behaviour across types (in particular but not limited to structs and enums). They are similar to interfaces in Java and Go.

```
trait Speak { // provide signature OR default method
    fn speak(&self) -> String; // Signature

    // fn speak(&self) -> String { // Default method
    //     "(silence)...".to_string()
    // }
}
```

```
// Implement the trait for Dog
struct Dog;
impl Speak for Dog {
    fn speak(&self) -> String {
        "Woof!".to_string()
    }
}

// Implement the trait for Cat
struct Cat;
impl Speak for Cat {
    fn speak(&self) -> String {
        "Meow!".to_string()
    }
}

// Generic function for ANY type implementing Speak
fn animal_talk<T: Speak>(animal: T) {
    println!("{}", animal.speak());
}

fn main() {
    let d = Dog;
    let c = Cat;

    animal_talk(d); // Woof!
    animal_talk(c); // Meow!
}
```

Standard traits and the derive macro

- Rust has many standard traits that you can implement. For example, you can implement the Add trait, to get operator overloading. The big four ...

```
// Display - User-facing string representation
impl fmt::Display for Point {
    fn fmt(&self, f: &mut fmt::Formatter) -> fmt::Result {
        write!(f, "({}, {})", self.x, self.y)
    }
}

// Debug - Developer-facing (and usually derived)
#[derive(Debug)] // Gives you {:?} formatting

// Default - Provides default values
impl Default for Config {
    fn default() -> Self {
        Config { timeout: 30, retries: 3 }
    }
}
let cfg = Config::default();

// From/Into - Type conversions
impl From<i32> for Number {
    fn from(item: i32) -> Self {
        Number { value: item }
    }
}
let num = Number::from(42);
let num: Number = 42.into(); // Into comes with From
```

- Derive macros automatically implement standard traits for your types, reducing boilerplate code.

```
#[derive(Debug, Clone, Copy, PartialEq, Eq)]
struct Point {
    x: i64,
    y: i64,
}
// Clone - enables explicit duplication with .clone()
// Copy - enables implicit copying (stack-only types)
// Debug - enables formatting with {:?} and {:#?}
// Default - provide fn default() for a default value
// Eq - requires PartialEq - rigorous/total equality
// Hash - enables hashing
// Ord - rigorous comparison and ordering
// PartialEq - enables == and != operators
// PartialOrd - enables comparison where possible
```


Generics

- Generics let you write a function, struct, enum or trait once, and use it with different types.

```
// struct example
struct Point<T> {
    x: T,
    y: T,
}
let int_point = Point { x: 5, y: -5 };
let float_point = Point{ x: 1.2, y: -2.1 };

// function example - with trait binding
fn biggest<T: PartialOrd + Copy>(list: &[T])
-> Option<T> {

    if list.is_empty() {
        return None;
    }
    let mut largest = list[0];
    for &item in list.iter() {
        if item > largest {
            largest = item;
        }
    }
    Some(largest)
}
```

Extending existing types in Rust

- The Problem:** You can't add methods directly to types you don't own (like Vec, String, etc.) due to Rust's orphan rule (which says, you can only implement a trait for a type if you own either the trait OR the type (or both)).
- Thin wrapper types** - give you full control to write your own implementation blocks, but you lose automatic type coercion. Also a bit clunky.

```
struct MyVec<T>(Vec<T>); // generic tuple struct

impl<T> MyVec<T> {
    fn double_len(&self) -> usize { self.0.len() * 2 }
    fn push(&mut self, item: T) { self.0.push(item); }
}
```

- Extension traits** allow you to write a new implementation block for an existing type. But to use the trait you need an explicit "use" statement.

```
trait StringExt {
    fn is_palindrome(&self) -> bool;
}

impl StringExt for String {
    fn is_palindrome(&self) -> bool {
        self == &self.chars().rev().collect::<String>()
    }
}

// Usage: must import trait / can be in same file
use StringExt;
let s = String::from("racecar");
println!("{}", s.is_palindrome()); // true
```

Smart pointers

- Smart pointers in Rust are data structures that act like pointers but add extra features for memory management, ownership, mutability, or thread safety.
- Box<T> is the simplest smart pointer. It puts a value on the heap instead of the stack. It is automatically freed when it goes out of scope.

```
let b = Box::new(5);
println!("{}", *b); // dereference to get 5
```

- Note: in the above example, although we can dereference b (because it implements the Deref trait), it is not really a reference – it is Box<i32> object that is owned by b.

- Rc<T> is reference counted, allowing multiple ownership. By keeping count of how many Rcs point to the same value. The value is freed when the last Rc goes out of scope.

```
use std::rc::Rc;

let a = Rc::new(String::from("hello"));
let b = Rc::clone(&a); // increase reference count
```

- Arc<T> is like Rc<T>, but thread safe because it uses atomic operations (that happen all at once).
- Related objects: Cell<T> and RefCell<T> which allows for the run-time mutability of immutable objects. Mutex<T> ensures only one thread can access the data at a time. RwLock<T> allows multiple readers or one writer.

Modules, Libraries and Crates

Modules

- A module is Rust's way to organize code into namespaces and control visibility (privacy). It allows one to group related functionality together.
- Declaration:** modules can be declared in three ways
 - As an inline module code-block
 - Every file.rs is a module
 - Every directory with mod.rs is a module

```
mod my_module {...} // Inline module named my_module
mod utils; // File-based (utils.rs)
mod network; // Directory-based (network/mod.rs)
```

Visibility

```
pub fn public() {} // Public to parent
fn private() {} // Private (default)
pub(crate) fn crate_only() {} // Visible within crate
pub(super) fn parent_only() {} // Visible to parent only
```

Use

```
my_module::function(); // Full path
use my_module::function; // Import single item
use my_module::{a, b, c} // Import multiple items
use my_module::*; // Import all public (avoid)
use super::parent_item; // From parent module
use crate::root_item; // From crate root
```

File structure

```
src/
├─ main.rs (and/or lib.rs) // Crate/Project root
├─ utils.rs // Simple module
└─ network/ // Directory module
    ├─ mod.rs // Module entry point
    └─ client.rs // Submodule
```

- Note: main.rs is used for executables that contain a main() function, and the lib.rs filename is used as the root file for libraries that do not contain a main() function. Both might be present when you have an executable that uses the library code.

Libraries and external crates

- Rust has a standard library that provides extra functionality, and there are many publicly available crates (or libraries) of code on the crates.io website.
- You import the standard library and external crates with the "use" keyword. For example:

```
// for collections
use std::collections::{HashMap, HashSet, BTreeMap};

// for IO and the file system
use std::fs;
use std::io::{self, Read, Write};
use std::path::Path;

// for networking
use std::net::{TcpListener, TcpStream, UdpSocket};
```

- To use library packages from the crates.io website, you will need to first identify the package names in a cargo.toml file in the root directory of your project, under a [dependencies] heading.
- For a project that uses the random number package, you might have a cargo.toml file like this.

```
[package]
name = "rand-example"
version = "0.1.0"
edition = "2024" # latest rust - released Feb 2025

[dependencies]
rand = "0.8"
```

- Inside your code it might look like this.

```
use rand::{thread_rng, Rng};
use rand::seq::SliceRandom; // for .choose()

fn main() {
    let mut rng = thread_rng();
    println!("Dice: {}", rng.gen_range(1..=6));
    println!("Float: {:.2}", rng.gen::());
    println!("Color: {}",
        ["R", "B", "G"].choose(&mut rng).unwrap());
}
```

- Running cargo build or cargo run will automatically download, compile, and link the dependency.

Testing

Basics

- Rust has built-in support for testing with the #[test] attribute. You write test functions inside a #[cfg(test)] module, and run them with cargo test.

```
pub fn add(a: i32, b: i32) -> i32 {
    a + b
}

#[cfg(test)]
mod tests {
    use super::*;

    #[test]
    fn adds_two_numbers() {
        assert_eq!(add(2, 3), 5);
    }

    #[test]
    #[should_panic(expected = "specific message")]
    fn test_specific_panic() {
        panic!("specific message");
    }
}
```

- The assert! family of macros are the backbone of testing in Rust. Their role is simple: they check that some condition holds true during a test, and if it doesn't, the test fails.
- Note: you can have more than one mod tests (or any number of #[cfg(test)] modules) in a single Rust file.
- Unit tests should appear in the same file as the functions/methods being tested.
- Integration tests should live in a tests/ subdirectory in the src directory.

- Tests are run from the command line using cargo.

```
cargo test
```

Compiler attributes

Attributes

- Attributes are code annotations that instruct the compiler, tools, or runtime about how to handle your code. They're written with a #[...] or #![...] syntax, and are essential for many Rust features.

- Examples follow:

```
#[allow(dead_code)] // suppress unused function warning
fn unused() {}

#[cfg(target_os = "linux")] // Conditional compilation
fn run_on_linux() {}

#[derive(Debug, Clone, PartialEq)] // Derive traits
struct Point { x: i32, y: i32 }

// This next one is a crate-level annotation
#![deny(warnings)] // fail build on any warning
```

Built in attributes you will see often

```
#[derive(...)] // Auto-implemented traits
#[cfg(...)] // Conditional compilation
#[test] // Mark test functions
#[allow(warning)] // Suppress specific warnings
#[repr(...)] // Control memory layout
#[inline] // Inlining hints
#[must_use] // Warn about unused results
#[deprecated] // Mark as deprecated
```

Concurrency / Parallelism

Some definitions

- **Parallelism** means your code is running on multiple CPUs at the same time (it's about execution). The standard threads library and the third-party Rayon crate offer strong parallelism support. Best for mathematical computations, data processing, CPU-heavy algorithms where you can divide work across multiple cores.
- **Concurrency** means dealing with multiple things at the same time whether that is happening on one CPU with context switching, or multiple CPUs (or perhaps many tasks running on just a few CPUs with some context switching). It's about program design, without necessarily a commitment to parallelism. Tokio and smol are two crates that support concurrency. Best for managing tasks that can spend a lot of time waiting for I/O.

Parallelism – Threads and Communication

- Basic thread spawning

```
use std::thread;
use std::time::Duration;

let handle = thread::spawn(|| {
    println!("Hello from thread!");
    // simulate work ...
    thread::sleep(Duration::from_millis(100));
    42 // return value
});
let result = handle.join().unwrap(); // wait/get result
```

- The `.join()` method in threading is used to block the calling thread until the thread on which `.join()` is called completes execution.

- Thread communication via channels

```
use std::sync::mpsc; // multi-producer, single-consumer

let (tx, rx) = mpsc::channel();
let tx2 = tx.clone(); // multiple senders allowed

thread::spawn(move || {
    tx.send("Hello").unwrap(); // move tx into thread
});
thread::spawn(move || {
    tx2.send("World").unwrap();
});

// Receive messages (blocks until available)
println!("{}", rx.recv().unwrap());
println!("{}", rx.recv().unwrap());
```

- Shared state with `Arc<Mutex<T>>`

```
use std::sync::{Arc, Mutex};

let counter = Arc::new(Mutex::new(0));
let counter2 = Arc::clone(&counter);

thread::spawn(move || {
    let mut num = counter2.lock().unwrap();
    *num += 1; // safely modify shared data
});
```

Concurrency using Tokio, async and await

```
// Async/await example (requires tokio or async-std)
use tokio::time::{sleep, Duration};

async fn fetch_data() -> String {
    sleep(Duration::from_millis(100)).await;
    "Data ready".to_string()
}

#[tokio::main] // macro creates runtime
async fn main() {
    let result = fetch_data().await; // wait until ready
    println!("{}", result);

    // Run multiple tasks concurrently
    let (r1, r2) = tokio::join!(
        fetch_data(),
        fetch_data()
    ); // both run in parallel
    println!("{}", r1 and r2);
}
```

What are Trait Objects?

- Trait objects (dyn Trait) enable runtime polymorphism, allowing different types to be used interchangeably through dynamic dispatch.

Set-up

```
trait Animal {
    fn speak(&self);
}

struct Dog;
struct Cat;

impl Animal for Dog {
    fn speak(&self) { println!("Woof!"); }
}

impl Animal for Cat {
    fn speak(&self) { println!("Meow!"); }
}

// Decide at compile time which method to use
fn static_sound<T: Animal>(animal: &T) {
    animal.speak();
}

// Decide at run-time which method to use
fn dynamic_sound(animal: &dyn Animal) {
    animal.speak()
}

fn main() {
    let dog = Dog;
    let cat = Cat;

    static_sound(&dog); // Woof!
    dynamic_sound(&cat); // Meow!
}
```

Object-Safe Traits necessary for dynamic dispatch

- Traits must be object-safe for dyn Trait. Cannot have:
 - Generic methods
 - Methods returning Self
 - Static methods without self

Which one to use:

- Static: faster, but duplicate copies of the function, or
- Dynamic: slightly slower, but more flexible.

Unsafe Rust

Rationale

- Unsafe Rust allows you to do things that the compiler cannot verify as safe, and you take on the responsibility for ensuring your code upholds Rust's memory safety guarantees.
- Here we are manually managing memory that Rust can't track (uninitialized vector elements).

```
use std::ptr;
let mut vec = Vec::with_capacity(3);
unsafe {
    // Write directly to uninitialized memory
    ptr::write(vec.as_mut_ptr(), 42);
    ptr::write(vec.as_mut_ptr().add(1), 84);
    vec.set_len(2); // Fix length after manual init
}
println!("{}", vec); // [42, 84]
```